

University of Waterloo
Department of Computer Science

CS370 Midterm Examination: Fall 2006

Thursday, November 2, 2006
Duration = 2 hours

Instructors: Y. Li

Name

Student ID

Section: 8:30 11:30

The aids allowed are:

- Printed Course Notes
- Lecture notes
- Hand Calculators

There are 7 questions - do all 7.

Question	Mark	Max	Init.
1		8	
2		6	
3		5	
4		8	
5		6	
6		8	
7		9	
Total		50	

1. (8 marks) Answer the following questions and briefly explain. Each of the four questions is worth 2 marks. **Note that you only obtain 2 marks if the answer is correct and the explanation is clear.**

- (a) **True or False.** If A is a n -by- n nonsingular matrix, then $\text{cond}(A) = \text{cond}(A^{-1})$. **Briefly explain.**

Solution: Statement is true. Recall that:

$$\begin{aligned}\text{cond}(A) &= \|A\| \|A^{-1}\| \\ \text{cond}(A^{-1}) &= \|A^{-1}\| \|(A^{-1})^{-1}\| \\ &= \|A^{-1}\| \|A\| \quad \text{since } (A^{-1})^{-1} = A\end{aligned}$$

Since multiplication is commutative, we have that $\text{cond}(A) = \text{cond}(A^{-1})$.

- (b) **True or False.** In solving a nonsingular system of linear equations, Gaussian elimination with partial pivoting usually yields a small residual even if the matrix is ill-conditioned. **Briefly explain.**

Solution: Statement is true. The computed solution \hat{x} from Gaussian elimination with partial pivoting satisfies

$$(A + \delta A)\hat{x} = b, \quad \text{where } \|\delta A\| \leq \rho \|A\| \epsilon_{\text{mach}}$$

and ρ is rarely greater than 10 and ϵ_{mach} is the machine epsilon. Thus the residual $r = A\hat{x} - b$ satisfies

$$\|r\| \leq \|\delta A\| \cdot \|\hat{x}\| \leq \rho \cdot \|A\| \cdot \|\hat{x}\| \cdot \epsilon_{\text{mach}}$$

or

$$\frac{\|r\|}{\|\hat{x}\| \|A\|} \leq \rho \cdot \epsilon_{\text{mach}}$$

Hence the statement is true.

(c) It is known that google pagerank is the unique solution to

$$Ax = x, \quad \sum_{i=1}^n x_i = 1, \quad \text{where } A \text{ is a google matrix.}$$

True or False. The google PageRank can mathematically be determined by solving the least squares problem

$$\begin{pmatrix} A - I \\ e^T \end{pmatrix} x = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$

Here e is the n -column vector of all ones and 0 is the n -vector of all zeros.

If the answer is no, **briefly explain** why not.

If yes, is computing the PageRank x by solving the least squares problem above a good computational approach when the total number of web pages is very large? **Explain.**

Solution: Statement is true. However, solving the least squares problem above (using either a normal equation approach or a QR factorization method) isn't a very good approach since it requires forming the coefficient matrix

$$\begin{pmatrix} A - I \\ e^T \end{pmatrix}$$

which is now completely dense even when the original connectivity matrix G is very sparse. Note that the google matrix A has strictly positive elements and is completely dense. We can no longer take advantage of the sparsity of the connectivity matrix G .

- (d) Consider the following two Matlab computations for a power iteration for Google pagerank. Let $\hat{G} = p * G * D$.

$$(M1) \quad (\hat{G} + e * z^T) * x, \quad (M2) \quad \hat{G} * x + e * (z^T * x)$$

Assume that $\hat{G}x$ requires $2(1 - \text{sparsity}(G))n^2$ floating point operations (flops), where $\text{sparsity}(G)$ is a percentage number representing the percentage of zero entries in G . Which statement below is true?

- (1) Computation (M1) requires more flops than (M2).
- (2) Computation (M2) requires more flops than (M1).
- (3) Computation (M1) and (M2) require the same number of flops.

Explain your answer by providing the number of flops required in each case.

Solution: This question requires detailed analysis of flops for each method.

In general, (M1) and (M2) require different number of flops. Statement (1) is true when n is large.

Computation (M1) requires:

- n^2 flops for computing $R1 = e * z^T$ (assuming no exploitation of the fact that components of e are ones)
- n^2 flops for $R2 = \hat{G} + R1$
- $n * (2 * n - 1)$ flops for computing $R2 * x$

Thus (M1) requires a total of $4n^2 - n$ flops.

Computation (M2) requires:

- $2n - 1$ flops for computing $\alpha = z^T x$
- n flops for computing $\alpha * e$
- $2(1 - \text{sparsity}(\hat{G}))n^2$ flops for computing $\hat{G} * x$

Thus (M2) requires $2(1 - \text{sparsity}(\hat{G}))n^2 + 3n - 1$ flops. The dominant cost of (M2) is $2(1 - \text{sparsity}(\hat{G}))n^2$, which is always smaller than the dominant cost $4n^2$ of (M1).

2. (6 marks) Consider the floating point number system $F(10, 5, -5, 10)$ with rounding to the nearest rule.

(a) What is the unit rounding error?

Solution: Denote the unit rounding error as E :

$$E = \frac{1}{2} \times 10^{-4} = \frac{1}{2 \times 10^4} = 5 \times 10^{-5}$$

(b) Is the number $1/3$ in this FPNS? What is $\text{fl}(1/3)$?

Solution: The exact representation of $1/3$ is not included in this FPNS since

$$\text{fl}(1/3) = 0.33333 \times 10^0$$

(c) Show that floating point add $10.000 \oplus 0.0008$ has a relative error no greater than the unit rounding error.

Solution: Carry out the floating point add:

$$\begin{aligned} 10.000 \oplus 0.0008 &= 0.10000 \times 10^2 + 0.80000 \times 10^{-3} \\ &= 0.10001 \times 10^2 \end{aligned}$$

Determine the relative error which:

$$\begin{aligned} \text{Rel. Error} &= \frac{|10.0008 - 10.001|}{10.0008} \\ &= 1.998 \times 10^{-5} < E \end{aligned}$$

Clearly, the relative error is less than the unit rounding error found in (a).

3. (5 marks) Let x be any real number. Give an expression for the solution of the following problem

$$\min_x \left\| \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix} x - \begin{bmatrix} b_1 \\ \vdots \\ b_m \end{bmatrix} \right\|_2$$

Explain your answer.

Solution: We use the least squares fitting to solve the above problem. Note that least squares fitting can be applied since vectors are just matrices of size $m \times 1$. Therefore, we have an overdetermined system and least squares fitting can be used.

Define the following matrices:

$$A = \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix}; \quad A^t = [1 \quad 1 \quad \dots \quad 1]; \quad B = \begin{bmatrix} b_1 \\ \vdots \\ b_m \end{bmatrix}$$

Using least squares fitting, we get the following equations:

$$\begin{aligned} A^t A x &= A^t B \\ m x &= \sum_{i=1}^m b_i \\ x &= \frac{1}{m} \sum_{i=1}^m b_i \end{aligned}$$

Hence, the solution to the above problem is to take the average of all the b_i .

4. (8 marks) Find numerical values $S'(0)$ and $S'(3)$ for the cubic spline below

$$S(x) = \begin{cases} S_1(x) = 1 + \alpha x + \frac{2}{3}x^2 - \frac{1}{3}x^3 & x \in [0, 2] \\ S_2(x) = -7 + 13x - \frac{16}{3}x^2 + \beta x^3 & x \in [2, 3] \end{cases}$$

Solution: We first need to determine the values of α and β in $S(x)$. We will do this by solving the equations obtained from the properties of a cubic spline, i.e. continuity of $S(x)$ and $S'(x)$.

- First, $S(x)$ is assumed to be continuous on $[0, 3]$ which means:

$$\begin{aligned} S_1(2) &= S_2(2) \\ 1 + 2\alpha + \frac{8}{3} - \frac{8}{3} &= -7 + 26 - \frac{64}{3} + 8\beta \\ 1 + 2\alpha &= -\frac{7}{3} + 8\beta \\ \alpha &= 4\beta - \frac{5}{3} \end{aligned}$$

- The next equation is obtained based on the continuity of $S'(x)$ on $[0, 3]$.

$$S'(x) = \begin{cases} S'_1(x) = \alpha + \frac{4}{3}x - x^2 & x \in [0, 2] \\ S'_2(x) = 13 - \frac{32}{3}x + 3\beta x^2 & x \in [2, 3] \end{cases}$$

Considering $S'_1(2) = S'_2(2)$, we get:

$$\begin{aligned} S'_1(2) &= S'_2(2) \\ \alpha + \frac{8}{3} - 4 &= 13 - \frac{64}{3} + 12\beta \\ \alpha &= -7 + 12\beta \end{aligned}$$

- Combining the two equations for α , we get:

$$\begin{aligned} 4\beta - \frac{5}{3} &= -7 + 12\beta \\ \beta &= \frac{2}{3} \end{aligned}$$

and

$$\alpha = 4 \times \frac{2}{3} - \frac{5}{3} = 1$$

Consequently, we find:

$$S'(x) = \begin{cases} S'_1(x) = 1 + \frac{4}{3}x - x^2 & x \in [0, 2] \\ S'_2(x) = 13 - \frac{32}{3}x + 2x^2 & x \in [2, 3] \end{cases}$$

which implies:

$$S'(0) = 1 \quad ; \quad S'(3) = 13 - 32 + 18 = -1$$

5. (6 marks) Assume that (x, y, z) , where x, y, z are each an m -vector, represents sample coordinates of points on a closed 3-dimensional parametric curve $(x(t), y(t), z(t))$. Note that $x_1 = x_m, y_1 = y_m, z_1 = z_m$. Using the index of the coordinates as the parameter t , write a Matlab segment to plot a smooth parametric curve (with a refining factor of 10) which interpolates the given data (x, y, z) using cubic splines. You can assume that the vectors x, y, z have been initialized for you.

Solution: Here is an example of the Matlab code segment:

```
m = length(x);

% Build parameter t
t = 1:m;

% Build splines
x_s = csape(t,x,'periodic');
y_s = csape(t,y,'periodic');
z_s = csape(t,z,'periodic');

% Refine t
t_ref = 1:0.1:m;

% Evaluate spline for refined t
x_ref = ppval(x_s,t_ref);
y_ref = ppval(y_s,t_ref);
z_ref = ppval(z_s,t_ref);

% Plot
plot3(x_ref, y_ref, z_ref);
```


6. (8 marks)

Find a 3-by-3 permutation matrix P and a lower triangular matrix L with unit diagonal and an upper triangular matrix U such that $PA = LU$ for

$$A = \begin{pmatrix} -2 & 8 & -28 \\ 1 & 8 & 6 \\ 4 & 8 & 24 \end{pmatrix}$$

Solution: Carry out the following row operations:

- Swap rows (1) and (3) to have largest entry in first column on top row:

$$\begin{pmatrix} 4 & 8 & 24 \\ 1 & 8 & 6 \\ -2 & 8 & -28 \end{pmatrix}$$

- We carry out the following row operation: $(3) + 1/2 * (1)$ and obtain:

$$\begin{pmatrix} 4 & 8 & 24 \\ 1 & 8 & 6 \\ 0 & 12 & -16 \end{pmatrix}$$

- Now, carry out the following row operation: $(2) - 1/4 * (1)$ and we obtain:

$$\begin{pmatrix} 4 & 8 & 24 \\ 0 & 6 & 0 \\ 0 & 12 & -16 \end{pmatrix}$$

- Next, swap rows (2) and (3):

$$\begin{pmatrix} 4 & 8 & 24 \\ 0 & 12 & -16 \\ 0 & 6 & 0 \end{pmatrix}$$

- Carry out the following row operation: $(3) - 1/2 * (2)$ and obtain:

$$\begin{pmatrix} 4 & 8 & 24 \\ 0 & 12 & -16 \\ 0 & 0 & 8 \end{pmatrix}$$

Consequently, we have:

$$U = \begin{pmatrix} 4 & 8 & 24 \\ 0 & 12 & -16 \\ 0 & 0 & 8 \end{pmatrix} \quad P = \begin{pmatrix} 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{pmatrix} \quad L = \begin{pmatrix} 1 & 0 & 0 \\ -1/2 & 1 & 0 \\ 1/4 & 1/2 & 1 \end{pmatrix}$$

7. (9 marks)

Assume that A , B , C and D are given n -by- n matrices and that A is nonsingular. Assume that f , g and h are given n -by-1 vectors.

- (a) Write a Matlab fragment that computes n -vectors x , y , and z so that the following equations hold:

$$\begin{array}{rcl} Ax + By + Cz & = & f \\ Ay + Dz & = & g \\ Az & = & h \end{array}$$

You may use the **lutx** and **forward** and **backsub** functions as in your assignment, see also appendix. **Efficiency matters.**

Solution: Here is an example of a Matlab code segment:

```
n = length(h);

[ L,U,p] = lutx(A);
w = forward(L,h(p));
z = backward(U,w);

b = g-D*z;
w_2 = forward(L,b(p));
y = backward(U,w_2);

d = f-C*z - B*y;
w_3 = forward(L,d(p));
x = backward(U,w_3);
```

- (b) How many flops are required in your computation of x , y , and z ? Just provide the dominant term (including an accurate coefficient).

Solution: Work for the following operations:

- LU factorization: $\frac{2}{3}n^3 + \mathcal{O}(n^2)$
- Forward and backward solve: $2n^2 + \mathcal{O}(n)$
- Matrix-vector multiply: $2n^2 - n$

In the code from part (a), the work can be calculated as follows:

$$\begin{aligned}\text{Nbr flops} &= (\text{LU fact.}) + 3 \times (\text{For+back solve}) + 3 \times (\text{Mat-vec mult.}) + 3 \times (\text{Sum of vectors}) \\ &= \frac{2}{3}n^3 + \mathcal{O}(n^2) + 3[2n^2 + \mathcal{O}(n)] + 3[2n^2 - n] + 3n \\ &= \frac{2}{3}n^3 + \mathcal{O}(n^2)\end{aligned}$$

Appendix

```

function pp = csape(x,y,conds,valconds)
%CSAPE Cubic spline interpolation with various end-conditions.
%
%   PP = CSAPE(X,Y)
%
%   returns the cubic spline interpolant (in ppform) to the given
%   data (X,Y) using Lagrange end-conditions (see default in table below).
%
%   PP = CSAPE(X,Y,CONDS) uses the end-conditions specified in CONDS, with
%   default values (which depend on the particular conditions).
%
%   CONDS may be a *string* whose first character matches one of the
%   following: 'complete' or 'clamped', 'not-a-knot', 'periodic',
%   'second', 'variational', with the following meanings:
%
%   'complete'      : match endslopes (as given in VALCONDS, with
%                      default as under *default*)
%   'not-a-knot'    : make spline C3 across first and last interior
%                      break (ignoring VALCONDS if given)
%   'periodic'      : match first and second derivatives at first data
%                      point with those at last data point
%                      (ignoring VALCONDS if given)
%   'second'        : match end second derivatives (as given in VALCONDS,
%                      with default [0 0], i.e., as in variational)
%   'variational'   : set end second derivatives equal to zero
%                      (ignoring VALCONDS if given)
%   The *default*   : match endslopes to the slope of the cubic that
%                      matches the first four data at the respective end.
%
%-----%

function output = spline(x,y,xx)
%SPLINE Cubic spline data interpolation.
%   YY = SPLINE(X,Y,XX) uses cubic spline interpolation to find YY, the values
%   of the underlying function Y at the points in the vector XX. The vector X
%   specifies the points at which the data Y is given. If Y is a matrix, then
%   the data is taken to be vector-valued and interpolation is performed for
%   each column of Y and YY will be length(XX)-by-size(Y,2).
%
%   PP = SPLINE(X,Y) returns the piecewise polynomial form of the cubic spline
%   interpolant for later use with PPVAL and the spline utility UNMKPP.
%

```

```

% Ordinarily, the not-a-knot end conditions are used. However, if Y contains
% two more values than X has entries, then the first and last value in Y are
% used as the endslopes for the cubic spline. Namely:
%      f(X) = Y(:,2:end-1),    df(min(X)) = Y(:,1),    df(max(X)) = Y(:,end)

%-----%

function v=ppval(pp,xx)
%PPVAL Evaluate piecewise polynomial.
% V = PPVAL(PP,XX) returns the value at the points XX of the piecewise
% polynomial contained in PP, as constructed by SPLINE or the spline utility
% MKPP.
%
% See also SPLINE, MKPP, UNMKPP.

%-----%

%PLOT3 Plot lines and points in 3-D space.
% PLOT3() is a three-dimensional analogue of PLOT().
%
% PLOT3(x,y,z), where x, y and z are three vectors of the same length,
% plots a line in 3-space through the points whose coordinates are the
% elements of x, y and z.

%-----%

function [L,U,p] = lutx(A)
%LUTX Triangular factorization, textbook version
% [L,U,p] = lutx(A) produces a unit lower triangular matrix L,
% an upper triangular matrix U, and a permutation vector p,
% so that L*U = A(p,:)

%-----%

function x = forward(L,x)
% FORWARD. Forward elimination.
% For lower triangular L, x = forward(L,b) solves L*x = b.

%-----%

function x = backsubs(U,x)
% BACKSUBS. Back substitution.
% For upper triangular U, x = backsubs(U,b) solves U*x = b.

%-----%

```

Scrap paper

Scrap paper